

Agent Safety Posture — Four Defence-in-Depth Levels

User message / documents

(4) Injection scan + lifecycle middleware
Pattern scanner before prompt (OWASP LLM01)
HTTP 403 on write POST without confirmation

(1) Tool-schema boundary
17 MCP tools only — 9 read · 8 write
No arbitrary code · no unapproved API

No single layer bypasses all others — model output cannot override middleware

(2) Human confirmation gate
Explicit affirmative consent per write
confirmation turn ID → AGENT_ACTION_LOG

(3) Session key scope
EIP-7702: named tools · value cap · 24h TTL
Revocable · out-of-scope rejected on-chain

On-chain write · audit trail